NSBM Green University

Faculty of Computing

BSc (Hons) Data Science

DS403.3- Big Data Programming

Intermediate Report

Group - 2

Student ID	Student Name
24490	MRK Karunathilaka
24572	MKIM Rohana
24614	GAAS Ganegoda

Module Lecturer: Mr. Adhil Rushdy

Table of Contents

1. Batch Processing Implementation Steps

Step 01: Resource Group Creation

- **Action**: Created a resource group named **lambda-arch-rg** to centralize all project resources.

- **Purpose**: Ensures organized management and cost tracking.

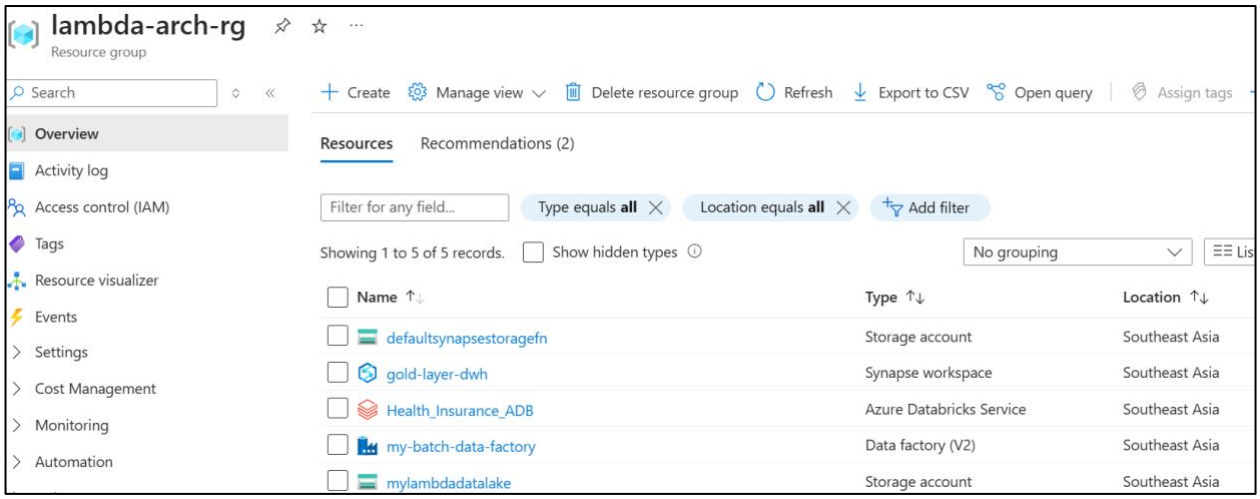

Step 02: Data Lake Setup

- **Actions**: Created a **Storage Account** (ADLS Gen2) with four containers,

    - **bronze:** Stores raw data (e.g., CSV files from Google Drive).
    - **silver:** Holds transformed/cleaned data (Parquet format).
    - **gold:** Stores analysis-ready datasets (aggregated tables).
    - **parameter:** Contains JSON files for dynamic pipeline configurations.

- **Purpose**: A Data Lake is very cost-effective; it can store both structured and unstructured data due to its object Storage.

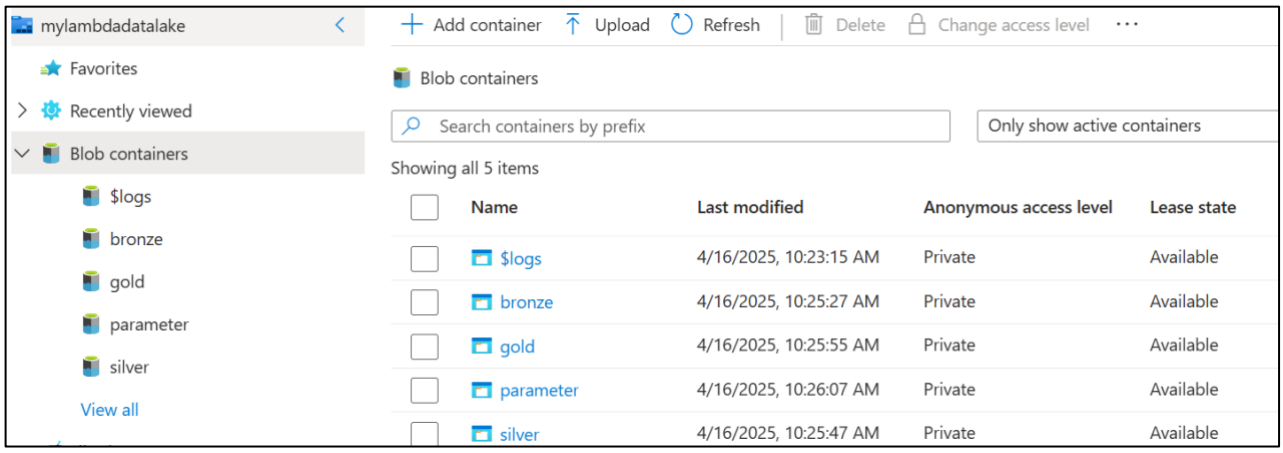

Step 03: Data Ingestion with Azure Data Factory (ADF)

- **Linked Services**
  - **ADLS Gen2**: Connected to the data lake containers.
  - **Google Drive**: Enabled CSV file ingestion (fallback after Git repo failed due to file size limits).

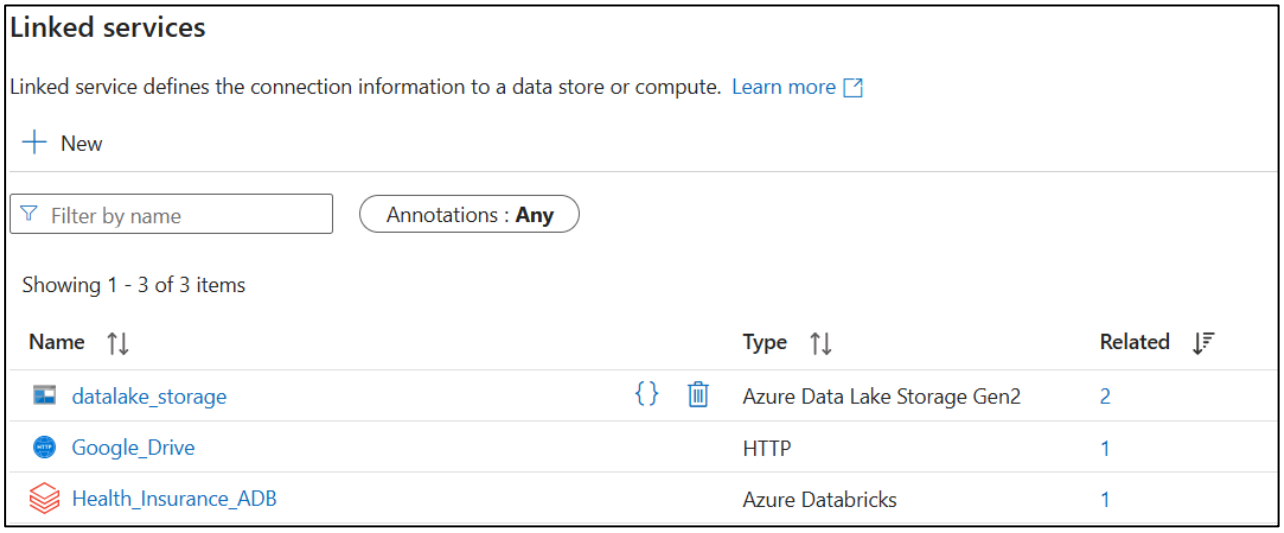

Challenges & Solutions

- **Google Drive 100MB Limit**
    - Split files into smaller chunks (<100MB) to avoid corruption.
- **Dynamic Pipeline**
    - Used **Lookup Activity** to fetch parameters from JSON files.
    - **ForEach Activity** + **Copy Activity** transferred files from Google Drive to bronze with,
        - Source parameter: relative_url.
        - Sink parameters: Folder_Name and File_Name.

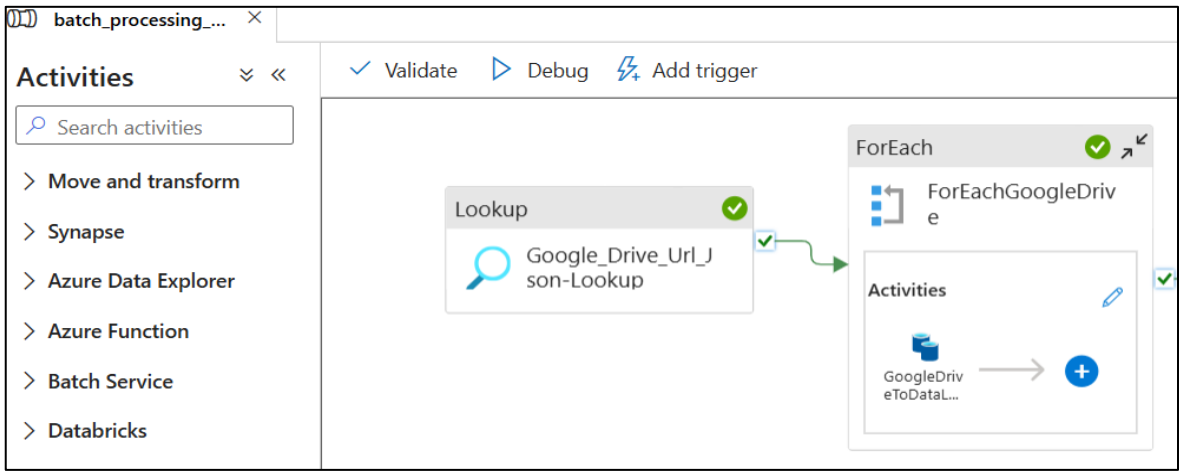

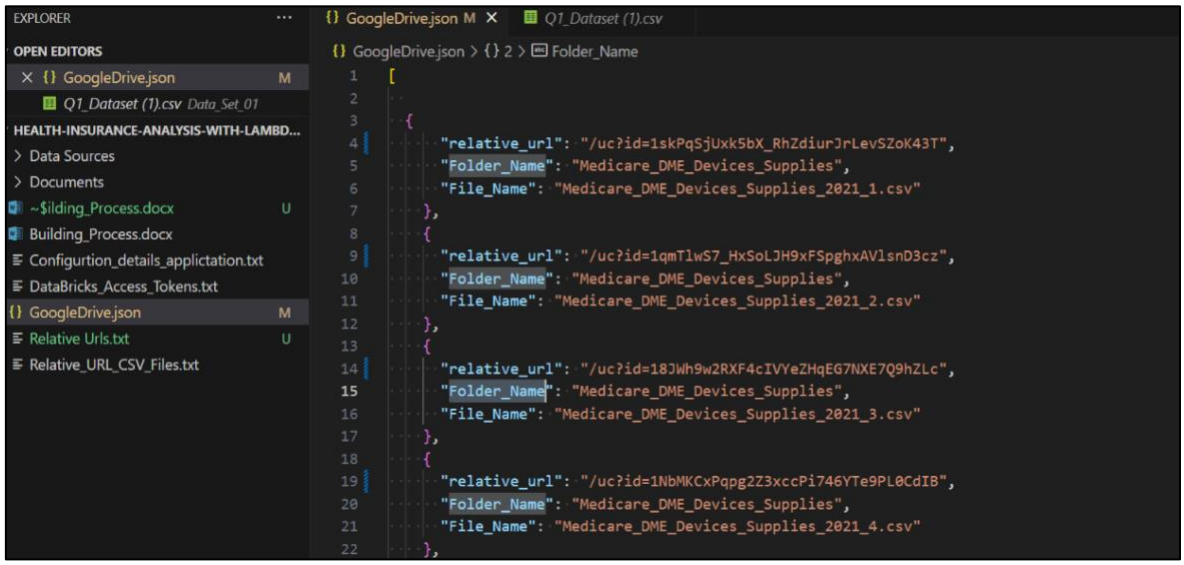

Step 04: Data Ingestion with Azure Data Factory (ADF)

- **Cluster Configuration**
    - **Single-Node Cluster** (Standard_DS3_v2, 14GB RAM) to minimize costs.
    - **Auto-termination**: Set to 10 minutes of inactivity.
- **Data Processing**
    1. Mounted **bronze** and **silver** containers to Databricks.
    2. Loaded data into Data Frames,
        - Claims_df (claim details).
        - Drugs_df (prescription data).
        - Medicare_DME_DS_df (medical equipment records).
    3. Transformed data (cleaning) → Saved to silver (Parquet).
- **Integration**
    - Linked Databricks notebook to ADF's batch_processing_pipeline via a Databricks-linked service.

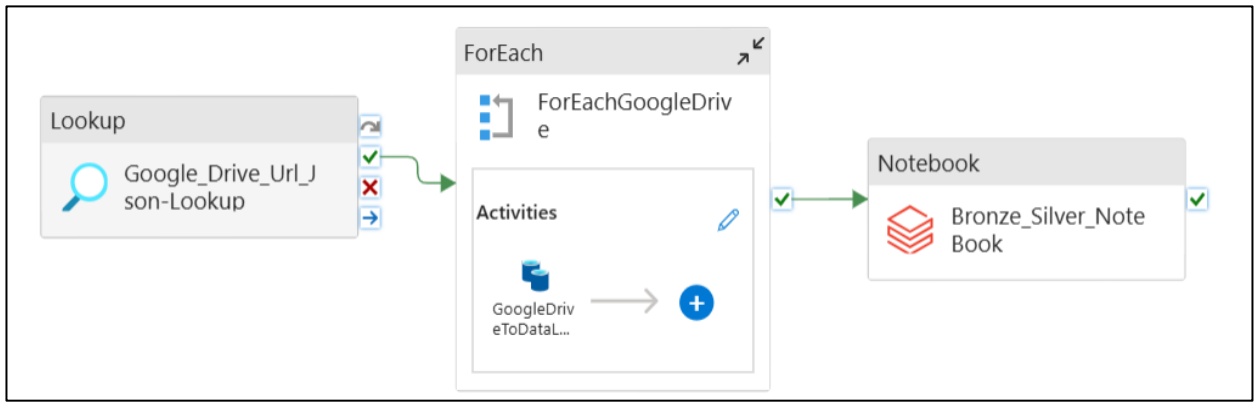

Step 05: Data Warehousing with Azure Synapse Analytics

- **Why Synapse?** Unified platform for,
    - **Data Factory (ADF)**: Pipeline orchestration (redundant with standalone ADF but retained for learning).
    - **Data Warehouse (DWH)**: Serverless SQL Pool chosen over Dedicated SQL Pool for:
        - **Cost Efficiency**: Pay-per-query (~$5/TB scanned) vs. fixed hourly costs.
        - **Data Virtualization**: Uses OPENROWSET() to query ADLS directly (no storage duplication).

2. Key Technical Decisions & Justifications in Cold Path

Component	Choice	Reason
Cluster Type	Single-Node (Databricks)	Cost savings; sufficient for batch workloads.
File Format	Parquet	Columnar storage → 80% smaller scans vs. CSV.
Synapse SQL Pool	Serverless	No infrastructure costs; scales to zero.
Data Ingestion	Google Drive + ADF	Workaround for Git's file size limits.

3. Challenges & Solutions in Cold Path

Challenge	Solution
Google Drive file corruption (>100MB)	Split files into sub-100MB chunks.
Databricks cluster startup delays	Auto-termination + single-node configuration.
Dynamic pipeline requirements	Parameterized JSON files + Lookup Activity.